

# **Kafka & Zookeeper Installation and other doc**

## Installation of Zookeeper

- Install java in the system
  - # yum install java -y
- Download the ZooKeeper from the link: <https://archive.apache.org/dist/zookeeper/zookeeper-3.5.7/apache-zookeeper-3.5.7-bin.tar.gz> using wget command
  - # wget <https://archive.apache.org/dist/zookeeper/zookeeper-3.5.7/apache-zookeeper-3.5.7-bin.tar.gz>
  - # tar -xf apache-zookeeper-3.5.7-bin.tar.gz
- We have config the configuration file for zookeeper, to do so go to the /kafka/apache-zookeeper-3.5.7-bin/conf/
  - # cp zoo\_sample.cfg zoo.cfg
  - Edit the zoo.cfg file with below code  
tickTime=2000  
initLimit=10  
syncLimit=5  
dataDir=/tmp/zookeeper  
clientPort=2181  
maxClientCnxns=60  
4lw.commands.whitelist=\*

```
[root@localhost conf]# ls
configuration.xml  log4j.properties  q  zoo.cfg  zoo_sample.cfg
[root@localhost conf]# cat zoo.cfg
tickTime=2000
initLimit=10
syncLimit=5
dataDir=/tmp/zookeeper
clientPort=2181
maxClientCnxns=60
4lw.commands.whitelist=*
```

- Install the nc and net-tools on using
  - # yum install nc -y
  - # yum install net-tools -y
- To start the zookeeper run below command and make sure you are in bin dir
  - # sh zkServer.sh start
- To check the status of zookeeper
  - # echo stat | nc localhost 2181

```
[root@localhost bin]# ls
README.txt  zkCleanup.sh  zkCli.cmd  zkCli.sh  zkEnv.cmd  zkEnv.sh  zkServer.cmd  zkServer-initialize.sh  zkServer.sh  zkTxnLogToolkit.cmd  zkTxnLogToolkit.sh
[root@localhost bin]# sh zkServer.sh start
/usr/bin/java
ZooKeeper JMX enabled by default
Using config: /kafka/apache-zookeeper-3.5.7-bin/bin/../conf/zoo.cfg
Starting zookeeper ... STARTED
[root@localhost bin]# echo stat | nc localhost 2181
Zookeeper version: 3.5.7-f0fdd52973d373ffd9c86b81d99842dc2c7f660e, built on 02/10/2020 11:30 GMT
Clients:
 /0:0:0:0:0:0:0:1:57488[0] (queued=0,recved=1,sent=0)

Latency min/avg/max: 0/0/0
Received: 1
Sent: 0
Connections: 1
Outstanding: 0
Zxid: 0x0
Mode: standalone
Node count: 5
```

# Installation of Kafka

- Download the kafka from the link: [https://downloads.apache.org/kafka/3.8.0/kafka\\_2.13-3.8.0.tgz](https://downloads.apache.org/kafka/3.8.0/kafka_2.13-3.8.0.tgz)
  - #wget [https://downloads.apache.org/kafka/3.8.0/kafka\\_2.13-3.8.0.tgz](https://downloads.apache.org/kafka/3.8.0/kafka_2.13-3.8.0.tgz)
  - #tar -xf kafka\_2.13-3.8.0.tgz

```
[root@localhost kafka_2.13-3.8.0]# cd bin
[root@localhost bin]# ls
connect-distributed.sh  kafka-cluster.sh      kafka-delete-records.sh  kafka-log-dirs.sh      kafka-run-class.sh      kafka-verifiable-consumer.sh  zookeeper-shell.sh
connect-mirror-maker.sh  kafka-configs.sh      kafka-dump-log.sh        kafka-metadata-quorum.sh  kafka-server-start.sh    kafka-verifiable-producer.sh
connect-plugin-path.sh   kafka-console-consumer.sh  kafka-e2e-latency.sh    kafka-metadata-shell.sh  kafka-server-stop.sh    trogdor.sh
connect-standalone.sh    kafka-console-producer.sh  kafka-features.sh        kafka-mirror-maker.sh    kafka-storage.sh         windows
kafka-acls.sh            kafka-consumer-groups.sh  kafka-get-offsets.sh     kafka-producer-perf-test.sh  kafka-streams-application-reset.sh  zookeeper-security-migration.sh
kafka-broker-api-versions.sh  kafka-consumer-perf-test.sh  kafka-jmx.sh            kafka-reassign-partitions.sh  kafka-topics.sh          zookeeper-server-start.sh
kafka-client-metrics.sh    kafka-delegation-tokens.sh  kafka-leader-election.sh  kafka-replica-verification.sh  kafka-transactions.sh     zookeeper-server-stop.sh
[root@localhost bin]# cd ..
[root@localhost kafka_2.13-3.8.0]# cd config/
[root@localhost config]# ls
connect-console-sink.properties  connect-file-sink.properties  connect-mirror-maker.properties  kraft  server.properties  zookeeper.properties
connect-console-source.properties  connect-file-source.properties  connect-standalone.properties  log4j.properties  tools-log4j.properties
connect-distributed.properties  connect-log4j.properties  consumer.properties  producer.properties  trogdor.conf
```

- Make change in the server.properties file

```
# The id of the broker. This must be set to a unique integer for each broker.
broker.id=1

##### Socket Server Settings #####

# The address the socket server listens on. If not configured, the host name will be equal to the value of
# java.net.InetAddress.getCanonicalHostName(), with PLAINTEXT listener name, and port 9092.
#   FORMAT:
#   listeners = listener_name://host_name:port
#   EXAMPLE:
#   listeners = PLAINTEXT://your.host.name:9092
listeners=PLAINTEXT://localhost:9092
```

- To start your kafka server in background run below
  - # sh kafka-server-start.sh -daemon /kafka/kafka\_2.13-3.8.0/config/server.properties
- To check kafka is up and running with brokers
  - # echo dump | nc localhost 2181 | grep brokers

```
[root@localhost config]# echo dump | nc localhost 2181 | grep brokers
/brokers/ids/1
[root@localhost config]#
```

- Logs of the kafka will be present in logs/server.log

```
[root@localhost logs]# ls
controller.log  kafka-authorizer.log  kafka-request.log  kafkaServer-gc.log  kafkaServer-gc.log.0  kafkaServer.out  log-cleaner.log  server.log  state-change.log
[root@localhost logs]# pwd
/kafka/kafka_2.13-3.8.0/logs
```

- Zookeeper logs will be available under /apache-zookeeper-3.5.7-bin/logs

```
[root@localhost logs]# ls
zookeeper-root-server-localhost.out
[root@localhost logs]# pwd
/kafka/apache-zookeeper-3.5.7-bin/logs
[root@localhost logs]#
```

## Creation of Topics in Kafka

- To create topic run below command where we create the topic with name “myTopic” with partitions as 1 and replication factor as 1

- #sh kafka-topics.sh --bootstrap-server localhost:9092 --create --topic myTopic --partitions 1 --replication-factor 1

```
[root@localhost bin]# sh kafka-topics.sh --bootstrap-server localhost:9092 --create --topic myTopic --partitions 1 --replication-factor 1
Created topic myTopic.
```

- To check the list of topic we have

- # sh kafka-topics.sh --bootstrap-server localhost:9092 --list

```
[root@localhost bin]# sh kafka-topics.sh --bootstrap-server localhost:9092 --list
myTopic
[root@localhost bin]#
```

- To describe the topic and check its setting

- # sh kafka-topics.sh --bootstrap-server localhost:9092 --describe --topic myTopic

```
[root@localhost bin]# sh kafka-topics.sh --bootstrap-server localhost:9092 --describe --topic myTopic
[2024-08-15 10:08:13,304] WARN [AdminClient clientId=adminclient-1] The DescribeTopicPartitions API is not supported, using Metadata API to describe topics. (org.apache.kafka.clients.admin.KafkaAdminClient)
Topic: myTopic TopicId: R62IuzcoREyQNeSIpokmVQ PartitionCount: 1 ReplicationFactor: 1 Configs:
Topic: myTopic Partition: 0 Leader: 1 Replicas: 1 Isr: 1 Elr: N/A LastKnownElr: N/A
```

- To publish/produce a message in the topic

- #sh kafka-console-producer.sh --bootstrap-server localhost:9092 --topic myTopic

```
[root@localhost bin]# sh kafka-console-producer.sh --bootstrap-server localhost:9092 --topic myTopic
>This is my first msg on kafka.
>
```

- To create the consumer to consume messages from the topic “myTopic”

- #sh kafka-console-consumer.sh --bootstrap-server localhost:9092 --topic myTopic --from-beginning

```
[root@localhost bin]# sh kafka-console-consumer.sh --bootstrap-server localhost:9092 --topic myTopic --from-beginning
This is my first msg on kafka.
```

- Whenever we don't specify the group for consumer, kafka will create group and tag consumer to it. To list down all the consumer groups

- # sh kafka-consumer-groups.sh --bootstrap-server localhost:9092 --list

- #sh kafka-consumer-groups.sh --bootstrap-server localhost:9092 --describe --group console-consumer-28586

```
[root@localhost bin]# sh kafka-consumer-groups.sh --bootstrap-server localhost:9092 --list
console-consumer-28586
[root@localhost bin]# sh kafka-consumer-groups.sh --bootstrap-server localhost:9092 --describe --group console-consumer-28586
```

GROUP	TOPIC	PARTITION	CURRENT-OFFSET	LOG-END-OFFSET	LAG	CONSUMER-ID	HOST	CLIENT-ID
console-consumer-28586	myTopic	0	-	2	-	console-consumer-5bfb239c-26b0-46a5-93c8-18c786d9254d	/127.0.0.1	console-consumer

- Kafka will manage all the details in \_\_consumer\_offsets where it will keep records how many messages are consumed by the consumer.

- # sh kafka-topics.sh --bootstrap-server localhost:9092 --list

```
[root@localhost bin]# sh kafka-topics.sh --bootstrap-server localhost:9092 --list
__consumer_offsets
myTopic
[root@localhost bin]#
```

- Multiple producers can produce messages on the topics and on the same way multiple consumers can consume messages from multiple topics

```
[root@localhost bin]# sh kafka-console-consumer.sh --bootstrap-server localhost:9092 -topic myTopic -group myConsumerGroup  
this is next messege
```

```
[root@localhost bin]# sh kafka-consumer-groups.sh --bootstrap-server localhost:9092 --list  
myConsumerGroup  
console-consumer-28586
```